

Submission Deadline

10th Mar 2024 (Monday) 11:59pm sharp. 2 marks penalty will be imposed on late submission (Late submission refers to submission or re-submission after the deadline). Submissions will be closed on **17th Mar 2024 (Monday) 11:59pm** sharp and no late submission will be accepted afterwards.

Objectives

In this assignment, you will implement a server that communicates with clients using an HTTP-like protocol. After completing this assignment, you should (1) be able to implement a simple TCP-based server, and (2) have a good understanding of the mechanisms of HTTP.

This programming assignment is worth 8 marks.

Group Work

All the work in this assignment should be done individually. However, if you find the assignment too difficult, you are allowed to form a group with another student (maximum two students per group). Group submission is subject to **2 marks penalty** for each student.

Under no circumstances should you solve it in a group and then submit it as an individual solution. This is considered plagiarism. There will be no acceptance for excuses such as forgetting to declare as group submission. Please refer to the "Special Instructions for Group Submission" and "Plagiarism Warning" on page 3 for more details.

Grading

We will test and grade your program on the `stu` server. Please make sure that your program run properly on `stu`. Moreover, you are allowed to use libraries installed in public folders of `stu` (e.g. `/usr/lib`) only.

We accept submission of Python 3 (in particular, 3.10), Java, or C/C++ program, and we recommend that you use **Python 3** for your assignments. Programming languages other than Python 3, Java, and C/C++ are not allowed. For Python 3, we use the `python3` program installed in folder `/usr/bin` on `stu` for grading. If you use Java or C/C++, we will compile and run your program for grading using the default compilers on `stu` (`openjdk 21` installed in `/usr/local/jdk/bin`, or `gcc 11.4.0` installed in `/usr/`

bin). The grading script infers your programming language from the file extension name (.py, .java, .c). Therefore, please ensure your files have the correct extension names. For a Java program, the class name should be consistent with the source file name, and please implement the static `main()` method so that the program can be executed as a standalone process after compilation. We will **deduct 1 mark** for every type of failure to follow instructions (e.g. wrong program name).

Note that for **Java** programs, name your main class as `WebServer` and hence source file name `WebServer.java` during development. When you want to test your program using the provided script on stu server, or to make submission, rename your file name according to program submission and group submission section **while keeping the class name as `WebServer`**. Grading script will rename your file accordingly during compilation.

We will grade your program based on its correctness only. A grading script will be used to test your program and no manual grading will be provided.

Testing Your Program

To test your program, please use your SoC UNIX ID and password to log on to `stu` as instructed on Assignment 0 paper. Your server program should receive one **command-line argument** which is the port number to `listen()`, as the following command shows:

```
python3 WebServer-A0165432X.py <port>
```

If you test your server manually, please select a port number greater than 1024, because the OS usually restricts usage of ports lower than 1024. If you get a `BindException: Address already in use` (or similar errors for other languages), please try a different port number. Note that your program should not read from `stdin`. Your program can print anything to `stdout` or `stderr`, and our test script will silently ignore them.

We also release a set of grading scripts to you under the **test** folder. **The scripts are intended for use on `stu` only**. There are no hidden test cases during grading. However, passing all the test cases does not guarantee that you will get full marks. We will detect fraudulent cases such as hard-coding answers.

To use the grading script, please upload your program **and all .so files** along with the `test` folder given in the package to `stu`. Make sure that your program, the .so files, and the `test` folder are in the same directory. Then, you can run the following command to test your server program:

```
bash test/WebServer.sh
```

By default, the script runs through all test cases. You can also choose to run a certain test case by specifying the case number in the command:

```
bash test/WebServer.sh 3
```

To stop a test, press `Ctrl-C`. If pressing the key combination once does not work, hold the keys until the script exits.

The test cases are embedded in the provided .so files, and therefore are not easily accessible. To check for the contents of a particular test case, use the provided query.sh file and specify the case number:

```
bash test/query.sh 3
```

There are four different test types in the test scripts. A brief description of each is as follows:

1. one connection per request - essentially the same as HTTP 1.0 (a new connection is opened each time a request has to be sent)
2. persistent connection, batched - persistent HTTP 1.1 with pipelining (all requests are sent as quickly as possible)
3. persistent connection, ping-pong - persistent HTTP 1.1 without pipelining (requests are sent one at a time, and the next is only sent upon receiving a complete response for the previous request)
4. intermittent connection - persistent HTTP 1.1, but requests can be split up into multiple chunks. There is no guarantee that a new request can only start at the beginning of a chunk, nor that a request can only end at the end of a chunk.

Requests should be responded to as soon as a complete request is received (see examples on page 8).

If you ever encounter this error: **tput: unknown terminal "xterm-256color"** when testing your program using script provided, run **export TERM=xterm** once after you log in and before you run WebServer.sh.

There is a low possibility that the grading script is grading an existing running server program instead of yours, if they use same port. So start the assignment early to avoid "congestion" during last few days. An advice to ensure the grading feedback you received is on your current program is to print to screen your student number right after receiving a connection, and kill your own existing server programs if any before testing it again. The grading result will be based on your program if you see your student number printed for particular test case, or otherwise try again.

Program Submission

For individual submission, please name your single source file as WebServer-<Matric number>.py and submit it to the Assignment 1 assignment on Canvas **without zip-ping**. Here, <Matric number> is your matriculation number which starts with letter A. An example file name would be WebServer-A0165432X.py. For resubmissions, keep the same file name before submitting to Canvas, and allow it to append numbers to the file. If you use Java, C, or C++ to implement the web server, please use .java, .c, or .cpp respectively as the extension name. Note that file names are **case-sensitive** on stu.

You are not allowed to post your solutions to any publicly accessible site on the

Internet.

Special Instructions for Group Submission

For group submission, please include matriculation numbers of both students in the file name, i.e. WebServer-<Matric number 1>-<Matric number 2>.py. Submit it to the same Assignment 1 assignment. An example file name would be WebServer-A0165432X-A0123456Y.py. For each group, there should be one designated member who submits the file, to avoid problems caused by multiple branches within a group. **Do not change the designated submitter!** If the group needs to upload a new version, it should be done by the same designated submitter as well.

Policy on AI Tools

Caution should be used when using AI tools (eg. ChatGPT). Usage of such tools may be considered plagiarism. Please see NUS' policy on AI for academic works:

<https://libguides.nus.edu.sg/new2nus/acadintegrity>

In general, the use of AI is acceptable so long as it does not involve code (eg. clarifying concepts taught during lecture, in English or other non-programming languages).

If AI is used for code (generation, or touching up of original code), the relevant sections of code should be marked accordingly with comments, but take note that excessive use of AI for code can be considered plagiarism in the same manner as use of code written by another person, as this assignment is meant to be your own work.

Example comments:

```
#AI Usage: Code generated by ChatGPT using prompt "reading standard input in Python"
(code)
```

```
#End section
```

or:

```
#AI Usage: Code touched up by ChatGPT using prompt "reading standard input in Python"
(code)
```

```
#End section
```

Any section of code not marked as such shall be considered a declaration that the code did not use AI for generation or touching up in any way.

Plagiarism Warning

You are free to discuss this assignment with your friends. However, you should refrain from sharing your program, program fragments, or detailed algorithms with others. If you want to solve this assignment in a group, please do so and **declare it as group work**.

We employ zero-tolerance policy against plagiarism. If a suspicious case is found, student would be asked to explain his/her code to the evaluator. A confirmed breach may result in an F grade for the entire course and further disciplinary action from the school.

Question & Answer

If you have any doubts on this assignment, please post your questions on Piazza before consulting the teaching team. However, the teaching team will NOT debug programs for students and we provide support for language-specific questions as a best-effort service. The intention of Q&A is to help clarify misconceptions or give you necessary directions.

The Server

In this assignment, you are required to implement a HTTP-like server that provides key-value store services. The server consists of two components. The first is a custom protocol simplified from HTTP 1.1. It retains the basic request-response mechanism and the persistent connection feature of HTTP 1.1, but uses a different header format. On top of this protocol, an interactive in-memory key-value store with a counter store are to be implemented under the path `/key/` and `/counter/` respectively. The key-value store supports insertion, update, retrieval, and deletion operations on key-value pairs. In addition, a record can be marked as temporary by adding a counter with the corresponding key in the counter store. The counter represents the remaining retrieval times of a temporary record in the key-value store. It will be decreased by 1 for each retrieval of the record in key-value store. A record in key-value store without a counter can be retrieved/updated with arbitrary times. In contrast, a temporary record is read-only and will be deleted after the corresponding counter reaches 0. Figure 1 shows the overall architecture of the web server.




Figure 1: The layered architecture of the server. You need to implement the top two layers which are highlighted with bold fonts.

The TCP Socket

The lowest-level layer of this server is the TCP socket interface. Basically, the server should

1. `bind()` to a port

2. `listen()` to the socket and `accept()` a new connection
3. Parse and respond to client request(s) sequentially, by using `recv()` and `send()` for socket I/O.
4. Going back to 2 after the client disconnects

As discussed in the "Testing Your Program" section, port number is passed to your server as a command line argument. The grading script ensures that only one client connects to your server at any time so that no connection multiplexing or polling is required; your server can simply process connections sequentially.

It is necessary to detect disconnection events reliably, because your server is required to support the persistent connection feature originated from HTTP 1.1. If the `bytes` object returned by `recv()` is of zero length, then the connection has closed. Note that if the connection is still open, but the client has no additional data to send (eg. client is waiting for a response), `recv()` may appear to stall until the connection is closed, instead of returning a `bytes` object. The client is guaranteed to keep the connection open until either the server finishes sending all responses or a timeout occurs. Therefore, there will be no `send()` failures on the server side, unless it times out during this process.

The HTTP-like Protocol

After establishing a new TCP connection with a client, your server should read a customized HTTP request from the socket, which consists of request headers and an optional request body. Your server parses the request and sends the corresponding response as defined below. After such a request-response cycle, the connection is **persistent** (i.e. not closing immediately), and the server waits for another request from the same client until a notification of socket disconnection.

Since a real-world TCP connection could be intermittent, your server is also required to work properly when requests are delivered in chunks of random sizes with delays between chunks. Note that a TCP connection only guarantees ordering of data, but not segment sizes or delays. The test script simulates this type of intermittent connection in some test cases, while keeping the delays between two consecutive transfers below one second. To handle chunks that contain incomplete requests, it is sufficient for your server to handle data buffering well.

In addition, your server should be **responsive** in the sense that upon receiving a complete request, the server processes it and sends back the response immediately. The test script sets a one-second timeout after each request is **fully** sent to the server. If a single chunk contains multiple requests, your server is also expected to respond quickly to all complete requests inside.

To grasp the core ideas of HTTP, we omit other tedious aspects of the protocol. For this assignment, the server only needs to handle basic request and response information, as well as the `Content-Length` header field. The **header format** of this HTTP-like

protocol is simplified as follows:

1. A header is a string consisting of **non-empty substrings** delimited with white-spaces (ASCII code 32). Two consecutive white-spaces mark the end of a header. You may assume that a substring does not contain white-spaces. Obviously, our header format is different from (and simpler than!) that of standard HTTP (e.g., no more `\r\n` or colons).
2. You may assume that a header consists of at least two substrings. The first two substrings (compulsory) contain information specific to requests or responses. Additional (and optional) substrings may follow. Every two substrings form a header field, where the first substring is the **case-insensitive** name of the header field, and the second is the value of this field.
3. For a request, the first two substrings (compulsory) are *HTTP method* and *path*, respectively. The HTTP method is **case-insensitive**, while path is **case-sensitive**. After these two substrings, optional header fields may follow, e.g. `Content-Length` and its value (a non-negative integer) if there is a content body in the request. The header finally ends with two white-spaces. This is followed by the content body, if any. Some sample requests are shown below:

- (a) `GET_/key/CS2105_`
- (b) `POST_/key/CS2105_Content-Length_27_Intro_To_Computer_Networks!`
- (c) `POST_/counter/CS2105_Content-Length_1_3`
- (d) `GET_/key/CS2105_`
- (e) `GET_/counter/CS2105_`
- (f) `POST_/counter/CS2105_Content-Length_1_2`
- (g) `DELETE_/key/CS2105_`
- (h) `DELETE_/counter/CS2105_`
- (i) `GET_/counter/CS2105_`
- (j) `DELETE_/key/CS2105_`

4. For a response, the first two substrings (compulsory) are *HTTP status code* and *description of status*. The description has **no** real effects and is kept as a convention as in HTTP, and it should not have any spaces. If the response has a content body, a `Content-Length` header should similarly be included. The corresponding changes of state in the server are shown in Figure 2 and responses for the above 10 requests could be:

- (a) `404_NotFound_`
- (b) `200_OK_`
- (c) `200_OK_`
- (d) `200_OK_Content-Length_27_Intro_To_Computer_Networks!`
- (e) `200_OK_Content-Length_1_2`
- (f) `200_OK_`

- (g) 405_MethodNotAllowed_
- (h) 200_OK_Content-Length_1_4
- (i) 200_OK_Content-Length_8_Infinity
- (j) 200_OK_Content-Length_27_Intro_To_Computer_Networks!




Figure 2: State of Key-value store and Counter store under the 9 consecutive requests.

The server can assume that **all requests are correctly formatted** as defined above. However, as in a real web server, a request may contain multiple header fields, and some of them may be unknown to your server (but with valid format). Therefore, Content-Length and its value, if they exist, may not be the first header field after the HTTP method and path, and it may not be the last as well. In addition, it is guaranteed that there are no duplicate header fields. Your server should parse all header fields and ignore irrelevant ones. Moreover, the maximum size of a content body is 5MB for this assignment.

The Key-value Store and Counter Store

A key-value store can be understood as a dictionary data structure that supports insertion, update, retrieval, and deletion of key-value pairs, with keys functioning as pointers to the values. If the key does not exist in the counter store, all four operations are applicable with arbitrary times e.g. (a), (b), (j). Otherwise, the key is marked as temporary and only the retrieval operation is allowed, e.g. (g). For each retrieval, the counter with the associated key should decrease by 1, e.g. (d). For simplicity, we restrict keys to be case-sensitive ASCII strings in this assignment. Due to the use of Content-Length header, values can safely be any **binary string**. The server should keep all data in memory only and avoid accessing the disk.

The counter store is very similar to key-value store, except that it stores positive remaining times for the retrieval, instead of bytes value, of a particular key in key-value store. It supports operations that are highly correlated with the key-value store: **Insertion** - which inserts a key to mark as temporary and the counter which indicates the number of retrieval times that this temporary record has remaining, e.g. (c), **Update** - which

increases the existing counter value by a certain positive integer, e.g. (f) **Retrieval** - which allows the client to retrieve the type of the key by getting it from the counter store which returns either the key is not found, positive counter of the temporary key, or a string value *Infinity* for a permanent key, e.g. (e) (i), and **Deletion** - which changes the temporary record back to permanent by deleting the counter in the counter store, e.g. (h). Counter keys are a subset of keys in key-value store i.e. only temporary keys should appear in the counter store. The content-length header and a positive integer content should be present in a POST.

Specifications

The key-value store should be accessible to the client through paths with prefix `/key/`. To operate on a specific key, the complete HTTP path is `/key/` appended with the actual key string. For example, a request to operate on key CS2105 specifies `/key/CS2105` as the HTTP path. The server should support the following key-value operations over the HTTP-like protocol:

1. Insertion and update

- (a) The HTTP request method should be `POST` and the value to be inserted (or updated) constitutes the content body. There should also be a `Content-Length` header indicating the number of bytes in the content body.
- (b) For insertions, the server should always respond with a `200 OK` status code after successfully inserting the key-value pair. For updates, one of two possibilities will occur.
 - If there is a positive number of retrieval times in the counter store for the associated key, then the update request is rejected and will return a `405 MethodNotAllowed` code.
 - Otherwise, server should always respond with a `200 OK` status code after updating the value.

2. Retrieval

- (a) The HTTP request method should be `GET`, and there is no content body.
- (b) One of three possibilities will occur.
 - If the key does not exist in the key-value store, the server returns a `404 NotFound` status code.
 - If the key already exists and there is a positive number of retrieval times in the counter store for the associated key, the server should return a `200 OK` code, the correct `Content-Length` header and the value data in the content body. Then, the server should decrease the remaining retrieval times of that key in the counter store by 1. Once the number of retrieval times reaches 0, the server should delete the key from both stores.

- Otherwise, if the key exists and is not in the counter store, the server should return a 200 OK code, the correct Content-Length header and the value data in the content body.

3. Deletion

- (a) The HTTP request method is DELETE, and there is no content body.
- (b) One of three possibilities will occur.
 - If the key does not exist, the server returns a 404 NotFound code.
 - If the key exists, but there is a positive number of retrieval times in the counter store for the associated key, then the delete request is rejected and will return a 405 MethodNotAllowed code.
 - Otherwise, it should delete the key-value pair from the store and respond with a 200 OK code and the deleted value string as the content body.

The counter store should be accessible to the client through paths with prefix `/counter/`. To operate on a specific key, the complete HTTP path is `/counter/` appended with the actual key string. For example, a request to operate on key CS2105 specifies `/counter/CS2105` as the HTTP path. Note the keys in counter store are a subset of the ones in key-value store, i.e. temporary keys should appear in both stores. The server should support the following counter key operations over the HTTP-like protocol for a counter store:

1. Insertion and Update

- (a) The HTTP request method should be POST with a positive integer in content body. The retrieval times of a corresponding key should be inserted (or incremented). For simplicity, you may assume that the testing script will only increment up to a maximum retrieval times of 9 for any key in the counter store.
- (b) For insertion, one of two possibilities will occur.
 - If the key does not exist in the key-value store, the server returns a 405 MethodNotAllowed code.
 - Otherwise, the server should respond with a 200 OK status code after updating the counter.
- (c) For updates, the server should increase the existing counter value by a positive integer specified in the request, and the client always expects success status.

2. Retrieval

- (a) The HTTP request method should be GET, and there is no content body.
- (b) One of three possibilities will occur.
 - If the key does not exist, and the key is not found in key-value store, the server returns a 404 NotFound code.

- If the key exists , the server should return a 200 OK code, the correct Content-Length header and the remaining retrieval times for the corresponding key.
- Otherwise, if the key is not found in the counter store, but exists in the key-value store, the server should return a 200 OK code, the correct Content-Length header and a string Infinity.

3. Deletion

- (a) The HTTP request method is DELETE, and there is no content body.
- (b) If the key does not exist, the server returns a 404 NotFound code. Otherwise, it should delete the key-value pair from the store and respond with a 200 OK code and the deleted counter integer as the content body. The Content-Length header should also be sent accordingly.

Grading Rubric

1. The server program is free from syntax errors and can be successfully executed (or compiled, for Java/C/C++). (1 mark)
2. Test cases (7 marks)
 - (a) You will receive marks based off the test cases passed.
 - (b) The starting number of marks is (-0.2) marks, before any test cases are considered. The actual marks awarded cannot fall below 0 marks.
 - (c) Each numbered test case from 1 to 18 is worth 0.4 marks, for a total of 7.2 marks.
 - (d) Therefore, in the event that you pass all test cases, you will receive 7 marks.
 - (e) For test case 1-17, each of test cases a, b, c are worth 0.1 mark, and an additional 0.1 mark is awarded for passing all 3.
 - (f) For test case 18, this is simply worth 0.4 marks.